

**NAME OF THE PROJECT**

# **Housing Price Prediction**

**Submitted by:**

Arshi Maitra

## Acknowledgement:

This project has been completed with the help of training documents and live classes recordings from Data Trained Education. Few helps on coding have also been taken from few data science websites like Toward Data Science, Geek for Geeks, Stack Overflow.

# INTRODUCTION

House price prediction can help the developer determine the selling price of a house and can help the customer to arrange the right time to purchase a house. There are three factors that influence the price of a house which include physical conditions, concept and location. Traditional house price prediction is based on cost and sale price comparison lacking of an accepted standard and a certification process. Therefore, the availability of a house price prediction model helps fill up an important information gap and improve the efficiency of the real estate market. Data science comes as a very important tool to solve problems in the domain to help the companies increase their overall revenue, profits, improving their marketing strategies and focusing on changing trends in house sales and purchases. Predictive modelling, Market mix modelling, recommendation systems are some of the machine learning techniques used for achieving the business goals for housing companies.

Below is the dataset available from the sale of houses in Australia as a US-based housing company named Surprise Housing has decided to enter the Australian market. The company uses data analytics to purchase houses at a price below their actual values and flip them at a higher price.

The main goal of the project will be to determine the following:

- To determine the price of houses in Australia
- To understand the factors that are responsible for the price of a house

# Analytical Problem Framing

Importing of necessary libraries:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import train_test_split
from sklearn.feature_selection import SelectKBest, f_classif, f_regression
from sklearn.linear_model import Ridge
from sklearn.linear_model import Lasso
from sklearn.metrics import r2_score
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import GridSearchCV
from prettytable import PrettyTable
import warnings
warnings.filterwarnings('ignore') #Importing the necessary Libraries
```

Data preprocessing/Data Cleaning:

Both the train and test set has been loaded

```
In [2]: Train_Data=pd.read_csv("HousingTrain.csv") #Checking the training dataset
```

```
In [3]: Test_Data=pd.read_csv("Housing Test.csv") #Checking the test dataset
```

```
In [4]: Train_Data.head()
```

```
Out[4]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature	MiscVal	Mo
0	127	120	RL	NaN	4928	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
1	889	20	RL	95.0	15865	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
2	793	60	RL	92.0	9920	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0	
3	110	20	RL	105.0	11751	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	MnPrv	NaN	0	
4	422	20	RL	NaN	16635	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0	

5 rows × 81 columns

## Going through the columns of both the data:

```
In [5]: Train_Data.columns #Checking the total columns
```

```
Out[5]: Index(['Id', 'MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea', 'Street',  
              'Alley', 'LotShape', 'LandContour', 'Utilities', 'LotConfig',  
              'LandSlope', 'Neighborhood', 'Condition1', 'Condition2', 'BldgType',  
              'HouseStyle', 'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd',  
              'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType',  
              'MasVnrArea', 'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual',  
              'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1',  
              'BsmtFinType2', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', 'Heating',  
              'HeatingQC', 'CentralAir', 'Electrical', '1stFlrSF', '2ndFlrSF',  
              'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath',  
              'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'KitchenQual',  
              'TotRmsAbvGrd', 'Functional', 'Fireplaces', 'FireplaceQu', 'GarageType',  
              'GarageYrBlt', 'GarageFinish', 'GarageCars', 'GarageArea', 'GarageQual',  
              'GarageCond', 'PavedDrive', 'WoodDeckSF', 'OpenPorchSF',  
              'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'PoolQC',  
              'Fence', 'MiscFeature', 'MiscVal', 'MoSold', 'YrSold', 'SaleType',  
              'SaleCondition', 'SalePrice'],  
              dtype='object')
```

```
In [6]: Test_Data.columns
```

```
Out[6]: Index(['Id', 'MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea', 'Street',  
              'Alley', 'LotShape', 'LandContour', 'Utilities', 'LotConfig',  
              'LandSlope', 'Neighborhood', 'Condition1', 'Condition2', 'BldgType',  
              'HouseStyle', 'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd',  
              'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType',  
              'MasVnrArea', 'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual',  
              'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1',  
              'BsmtFinType2', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', 'Heating',  
              'HeatingQC', 'CentralAir', 'Electrical', '1stFlrSF', '2ndFlrSF',  
              'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath',  
              'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'KitchenQual',  
              'TotRmsAbvGrd', 'Functional', 'Fireplaces', 'FireplaceQu', 'GarageType',  
              'GarageYrBlt', 'GarageFinish', 'GarageCars', 'GarageArea', 'GarageQual',  
              'GarageCond', 'PavedDrive', 'WoodDeckSF', 'OpenPorchSF',  
              'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'PoolQC',  
              'Fence', 'MiscFeature', 'MiscVal', 'MoSold', 'YrSold', 'SaleType',  
              'SaleCondition'],  
              dtype='object')
```

## Checking the null values from Train and Test data:

```
In [8]: Train_Data.isnull().sum() #Checking the null values of the train data
```

```
Out[8]: Id                0
        MSSubClass         0
        MSZoning            0
        LotFrontage       214
        LotArea            0
        Street            0
        Alley            1091
        LotShape          0
        LandContour       0
        Utilities         0
        LotConfig         0
        LandSlope         0
        Neighborhood      0
        Condition1        0
        Condition2        0
        BldgType          0
        HouseStyle        0
        OverallQual       0
```

```
LowQualFinSF            0
GrLivArea               0
BsmtFullBath            0
BsmtHalfBath            0
FullBath                0
HalfBath                0
BedroomAbvGr           0
KitchenAbvGr           0
KitchenQual            0
TotRmsAbvGrd           0
Functional              0
Fireplaces              0
FireplaceQu            551
GarageType              64
GarageYrBlt            64
GarageFinish            64
GarageCars              0
GarageArea              0
GarageQual              64
GarageCond              64
PavedDrive              0
WoodDeckSF              0
OpenPorchSF            0
EnclosedPorch           0
3SsnPorch              0
ScreenPorch             0
PoolArea                0
PoolQC                 1161
Fence                   931
```

```
OverallCond            0
YearBuilt               0
YearRemodAdd           0
RoofStyle               0
RoofMat1                0
Exterior1st            0
Exterior2nd            0
MasVnrType              7
MasVnrArea              7
ExterQual               0
ExterCond               0
Foundation              0
BsmtQual                30
BsmtCond                30
BsmtExposure            31
BsmtFinType1            30
BsmtFinSF1              0
BsmtFinType2            31
BsmtFinSF2              0
BsmtUnfSF               0
TotalBsmSF              0
Heating                 0
HeatingQC               0
CentralAir              0
Electrical              0
1stFlrSF                0
2ndFlrSF                0
```

MiscFeature	1124
MiscVal	0
MoSold	0
YrSold	0
SaleType	0
SaleCondition	0
SalePrice	0

dtype: int64

## Dropping the null values as well as filling up the null values:

```
In [9]: # For the train data- before imputing the null values, it can be seen that for few columns there are more 90% of null values.  
# Hence dropping those columns
```

```
In [10]: Train_Data.drop(['Alley', 'PoolQC', 'Fence', 'MiscFeature'], axis=1, inplace=True)
```

## Data Imputation:

```
In [13]: Train_Data['LotFrontage']=Train_Data['LotFrontage'].fillna(Train_Data['LotFrontage'].mean())  
Train_Data['BsmtQual']=Train_Data['BsmtQual'].fillna(Train_Data['BsmtQual'].mode()[0])  
Train_Data['BsmtCond']=Train_Data['BsmtQual'].fillna(Train_Data['BsmtCond'].mode()[0])  
Train_Data['BsmtExposure']=Train_Data['BsmtExposure'].fillna(Train_Data['BsmtExposure'].mode()[0])  
Train_Data['BsmtFinType1']=Train_Data['BsmtFinType1'].fillna(Train_Data['BsmtFinType1'].mode()[0])  
Train_Data['BsmtFinType2']=Train_Data['BsmtFinType2'].fillna(Train_Data['BsmtFinType2'].mode()[0])  
Train_Data['FireplaceQu']=Train_Data['FireplaceQu'].fillna(Train_Data['FireplaceQu'].mode()[0])  
Train_Data['GarageType']=Train_Data['GarageType'].fillna(Train_Data['GarageType'].mode()[0])  
Train_Data['GarageYrBlt']=Train_Data['GarageYrBlt'].fillna(Train_Data['GarageYrBlt'].mean())  
Train_Data['GarageFinish']=Train_Data['GarageFinish'].fillna(Train_Data['GarageFinish'].mode()[0])  
Train_Data['GarageCars']=Train_Data['GarageCars'].fillna(Train_Data['GarageCars'].mean())  
Train_Data['GarageArea']=Train_Data['GarageArea'].fillna(Train_Data['GarageArea'].mean())  
Train_Data['GarageQual']=Train_Data['GarageQual'].fillna(Train_Data['GarageQual'].mode()[0])  
Train_Data['GarageCond']=Train_Data['GarageCond'].fillna(Train_Data['GarageCond'].mode()[0])  
Train_Data['MasVnrType']=Train_Data['MasVnrType'].fillna(Train_Data['MasVnrType'].mode()[0])  
Train_Data['MasVnrArea']=Train_Data['MasVnrArea'].fillna(Train_Data['MasVnrArea'].mean())
```

## Encoding the object data types:

```
In [22]: from sklearn.preprocessing import LabelEncoder
```

```
In [23]: enc= LabelEncoder()
```

```
In [24]: columns=['MSZoning','Street','LotShape','LandContour','Utilities','LotConfig','LandSlope','Neighborhood','Condition1','Condition2']  
X[columns] = X[columns].apply(enc.fit_transform)
```

## Selecting K Best feature selection method:

As we had 81 columns hence it was difficult to train the model with all the 80 features. Therefore K Best feature selection method has been utilised in order to get the best 25 features in comparison to 80.

```
In [26]: Features= SelectKBest(score_func=f_regression, k=25)  
fit=Features.fit(X,Y)  
scores=pd.DataFrame(fit.scores_)  
columns=pd.DataFrame(X.columns)
```



## Checking the scores:

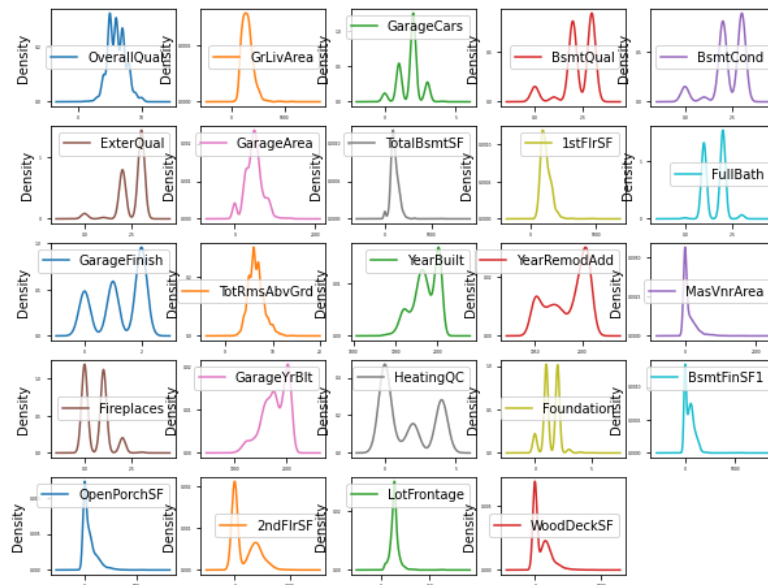
```
In [27]: Total_Score=pd.concat([columns,scores],axis=1)
Total_Score.columns=['Column','Score']
print(Total_Score.nlargest(25,'Score'))
```

	Column	Score
16	OverallQual	1925.310146
45	GrLivArea	1167.278131
60	GarageCars	760.625799
29	BsmtQual	754.737009
30	BsmtCond	754.737009
26	ExterQual	746.729944
61	GarageArea	724.282299
37	TotalBsmtSF	639.162482
52	KitchenQual	630.659595
42	1stFlrSF	615.032166
48	FullBath	518.998338
59	GarageFinish	472.788235
53	TotRmsAbvGrd	451.573798
18	YearBuilt	419.564317
19	YearRemodAdd	405.199690
25	MasVnrArea	319.254378
55	Fireplaces	312.273172
58	GarageYrBlt	309.521298
39	HeatingQC	230.953343
28	Foundation	189.817339
33	BsmtFinSF1	176.819524
66	OpenPorchSF	151.901231
43	2ndFlrSF	142.869824
3	LotFrontage	136.550260
65	WoodDeckSF	128.843420

Therefore now received the best 25 features for the target variable.

## Checking Data Skewness :

```
In [30]: NewX.plot(kind='kde',subplots=True,layout=(5,5),sharex=False,legend=True,fontsize=3,figsize=(10,8))
plt.show() # plotting the data and observing high skewness
```



## Skewness Score:

```
In [31]: NewX.skew().sort_values(ascending=False) #checking the skewness
```

```
Out[31]: MasVnrArea      2.834658
LotFrontage    2.710383
OpenPorchSF    2.410840
BsmtFinSF1     1.871606
TotalBsmtSF    1.744591
1stFlrSF       1.513707
WoodDeckSF     1.504929
GrLivArea      1.449952
2ndFlrSF       0.823479
Fireplaces     0.671966
TotRmsAbvGrd   0.644657
HeatingQC      0.449933
GarageArea     0.189665
OverallQual    0.175082
FullBath       0.057809
Foundation     -0.002761
GarageCars     -0.358556
GarageFinish   -0.450190
YearRemodAdd   -0.495864
YearBuilt      -0.579204
GarageYrBlt    -0.662934
BsmtCond       -1.343781
BsmtQual       -1.343781
ExterQual      -1.810843
dtype: float64
```

## Importing Power Transformation.

As the data is skewed it will affect the prediction score. Therefore data has to be transformed before taking any further action:

```
In [32]: from sklearn.preprocessing import power_transform
```

```
In [33]: New_X=power_transform(NewX)
```

```
In [34]: pd.DataFrame(New_X,columns=NewX.columns).skew().sort_values(ascending=False) # transforming the data to reduce skewness
```


```
Out[34]: MasVnrArea      0.416370
TotalBsmtSF    0.286779
2ndFlrSF       0.280208
LotFrontage    0.161368
HeatingQC      0.156511
WoodDeckSF     0.113026
Fireplaces     0.084950
OverallQual    0.021658
Foundation     0.004296
TotRmsAbvGrd   0.002332
GrLivArea     -0.000054
1stFlrSF      -0.002391
OpenPorchSF    -0.002749
GarageCars     -0.022970
FullBath       -0.045944
YearBuilt      -0.126641
GarageYrBlt    -0.132523
YearRemodAdd   -0.225131
GarageArea     -0.320370
GarageFinish   -0.335248
BsmtFinSF1     -0.404528
BsmtCond       -0.413999
BsmtQual       -0.413999
```

```
ExterQual      -0.605112
dtype: float64
```

## Checking Multicollinearity:

### i) Heat Map:

```
In [37]: corr_mat=X.corr()
plt.figure(figsize=[20,25])
sns.heatmap(corr_mat,annot=True)
plt.title("Correlation Matrix")
plt.show() #Checking correlation
```



From the heat map it can be observed that basement quality and basement condition has 100 % multicollinearity issue. Therefore it needs to be dropped, lets check via another metric that is VIF.

## ii) VIF:

```
In [39]: vif_data = pd.DataFrame()
vif_data["feature"] = X.columns
vif_data["VIF"] = [variance_inflation_factor(X.values, i)
                    for i in range(len(X.columns))]
vif_data
```

	feature	VIF
0	OverallQual	3.315234
1	GrLivArea	19.498638
2	GarageCars	5.351208
3	BsmtQual	inf
4	BsmtCond	inf
5	ExterQual	2.402698
6	GarageArea	4.981370
7	TotalBsmtSF	3.060318
8	1stFlrSF	12.212694
9	FullBath	2.358885
10	GarageFinish	1.946111
11	TotRmsAbvGrd	3.644552
12	YearBuilt	5.973074
13	YearRemodAdd	2.279980
14	MasVnrArea	1.385100
15	Fireplaces	1.529214
16	GarageYrBlt	4.344212
17	HeatingQC	1.652739
18	Foundation	2.007496
19	BsmtFinSF1	1.290198
20	OpenPorchSF	1.475894
21	2ndFlrSF	11.942552
22	LotFrontage	1.334319
23	WoodDeckSF	1.181661

As concluded from the heat map above, Basement Quality and basement condition has 100 % multicollinearity issue. Along with that first floor square feet, second floor squarefeet as well as Ground level above squarefeet also has multicollinearity issue.

It can be said be Ground level above squarefeet is the same data as first floor square feet, second floor squarefeet.

Therefore Basement Quality, first floor square feet and second floor squarefeet are dropped off in the train and test data.

```
In [40]: X.drop('BsmtQual',axis=1,inplace=True)
```

```
In [42]: X.drop(['1stFlrSF','2ndFlrSF'],axis=1,inplace=True)
```

### Checking the final VIF Score:

```
In [48]: vif_data = pd.DataFrame()
vif_data["feature"] = X.columns
vif_data["VIF"] = [variance_inflation_factor(X.values, i)
                    for i in range(len(X.columns))]
```

vif\_data

0	OverallQual	3.261709
1	GrLivArea	5.866728
2	GarageCars	5.347618
3	BsmtCond	2.148281
4	ExterQual	2.401001
5	GarageArea	4.962258
6	TotalBsmtSF	1.890722
7	FullBath	2.341211
8	GarageFinish	1.941324
9	TotRmsAbvGrd	3.641823
10	YearBuilt	5.935050
11	YearRemodAdd	2.274966
12	MasVnrArea	1.372453
13	Fireplaces	1.498093
14	GarageYrBlt	4.324579
15	HeatingQC	1.651327
16	Foundation	1.944356
17	BsmtFinSF1	1.289871
18	OpenPorchSF	1.465648
19	LotFrontage	1.295308
20	WoodDeckSF	1.177207

Therefore, the issue of multicollinearity has now been resolved.

### Detecting Outliers:

Zscore method has been applied to remove the outliers in the data.

```
In [51]: from scipy.stats import zscore
```

```
In [52]: (np.abs(zscore(X)<3)).all()
```

```
Out[52]: OverallQual      True
GrLivArea      False
GarageCars      False
BsmtCond        True
ExterQual        True
GarageArea      False
TotalBsmtSF     False
FullBath        True
GarageFinish     True
TotRmsAbvGrd    False
YearBuilt       True
YearRemodAdd    True
MasVnrArea      True
Fireplaces      True
GarageYrBlt     True
HeatingQC       True
Foundation      False
BsmtFinSF1      True
OpenPorchSF     True
LotFrontage     False
WoodDeckSF      True
dtype: bool
```

Outliers has been removed and values without outliers have been stored in a different variable.

```
In [53]: # these are the index positions where outlier is present
index = np.where(np.abs(zscore(X))>3)
index
```

```
Out[53]: (array([ 34,  48,  48,  48,  48,  52,  52,  54,  60,  86,  96,
119, 121, 124, 137, 141, 141, 159, 177, 191, 195, 195,
210, 211, 226, 231, 243, 249, 249, 249, 267, 305, 361,
361, 370, 420, 432, 483, 491, 498, 504, 510, 517, 537,
544, 558, 592, 592, 592, 592, 614, 644, 656, 691, 698,
706, 735, 747, 758, 760, 772, 800, 831, 834, 846, 846,
865, 884, 897, 899, 902, 908, 915, 935, 980, 980, 1025,
1035, 1042, 1046, 1053, 1056, 1067, 1094, 1104, 1104, 1107, 1117,
1120, 1126, 1144, 1147, 1148, 1164], dtype=int64),
array([ 6,  0,  1,  6,  9,  0,  1,  6,  6,  6,  6,  7,  2,  6,  6,  1, 19,
 6, 16, 19, 16, 19, 16,  7,  6, 19,  6,  0,  1,  7,  6,  6,  1,  6,
 6, 19,  6,  7,  6, 16,  6,  0,  6, 19,  7, 19,  1,  5,  6, 19,  9,
 7, 19,  1, 19, 19,  6, 19, 19, 19,  5, 19, 19,  6,  0,  7, 19,  6,
 5,  6, 19,  6, 19,  6,  2,  5,  2,  6,  6, 19, 19,  6,  6,  5,  2,
16,  6,  6, 19,  6, 16, 19,  6,  6], dtype=int64))
```

```
In [54]: New_X = X[(np.abs(zscore(X))<3).all(axis=1)]
New_X
```

```
In [55]: Y_new=Y.drop(index[0],axis=0)
Y_new #removing the outliers from target variables
```

```
Out[55]: 0      128000
1      268000
2      269790
3      190000
4      215000
...
1162    58500
1163    122000
1165    148500
1166    40000
1167    183200
Name: SalePrice, Length: 1089, dtype: int64
```

## Scaling the data:

```
In [60]: X_Scaled=Scalar.fit_transform(New_X)
X_Scaled
```

```
Out[60]: array([[ -0.09859076, -1.3363421,  0.25995446, ...,  1.38217617,
 0.10106507, -0.97594068],
 [ 1.39641172,  1.36740337,  0.25995446, ...,  1.38669106,
 1.27311012,  0.76580805],
 [ 0.65945083,  1.05642844,  0.25995446, ...,  1.17074992,
 1.13760085,  1.02781403],
 ...,
 [-0.09859076,  0.0126929,  0.25995446, ..., -1.0932726,
 -3.2731638,  0.78974467],
 [-1.69445354, -0.31064138, -1.12809175, ...,  0.76598274,
 -1.14710557, -0.97594068],
 [-0.09859076,  0.12790642,  0.25995446, ...,  0.91672675,
 0.10106507,  0.83612318]])
```

# **Model/s Development and Evaluation**

**(Linear Regression, Decision Tree Regressor,  
Random Forest Regressor and Gradient Boosting  
Regressor)**

**Linear Regression: 83.66%**

## Linear Regression

```
In [63]: LR=LinearRegression()
```

```
In [64]: X_train,X_test,y_train,y_test=train_test_split(X_Scaled,Y_new,test_size=0.20,random_state=50)
LR.fit(X_train,y_train)
pred_test=LR.predict(X_test)

print('R-Squared:',r2_score(y_test,pred_test)*100)
```

R-Squared: 83.66504953665496

**Decision Tree Regressor: 66.61%**

## Decision Tree Regressor

```
In [65]: DT=DecisionTreeRegressor()
```

```
In [66]: X_train,X_test,y_train,y_test=train_test_split(X_Scaled,Y_new,test_size=0.20,random_state=50)
DT.fit(X_train,y_train)
pred_test=DT.predict(X_test)

print('R-Squared:',r2_score(y_test,pred_test)*100)
```

R-Squared: 66.61352107885887

**Random Forest Regressor::85.52%**

## Random Forest Regressor

```
In [67]: rf=RandomForestRegressor()
```

```
In [68]: X_train,X_test,y_train,y_test=train_test_split(X_Scaled,Y_new,test_size=0.20,random_state=50)
rf.fit(X_train,y_train)
pred_test=rf.predict(X_test)

print('R-Squared:',r2_score(y_test,pred_test)*100)
```

R-Squared: 85.52395463859415

**Gradient Boosting Regressor:86.66%**

## Gradient Boosting

```
In [71]: GB=GradientBoostingRegressor()
```

```
In [74]: X_train,X_test,y_train,y_test=train_test_split(X_Scaled,Y_new,test_size=0.20,random_state=50)
GB.fit(X_train,y_train)
pred_test=GB.predict(X_test)

print('R-Squared:',r2_score(y_test,pred_test)*100)
```

R-Squared: 86.66187602268887

After getting the r squared scores for all the models, it needs to be checked whether any one of them are overfitted, hence cross validation technique has been used.



## Cross Validation for LR:

### Cross Validation for LR

```
In [76]: for i in range(2,6):  
         LR_Val=cross_val_score(LR,X_Scaled,Y_new,cv=i)  
         print("The cross validation score for Linear Regressor",i,"is",LR_Val.mean())
```

The cross validation score for Linear Regressor 2 is 0.8099555120612187  
The cross validation score for Linear Regressor 3 is 0.8165233974276257  
The cross validation score for Linear Regressor 4 is 0.8165572438206878  
The cross validation score for Linear Regressor 5 is 0.8132481644785304

## Cross Validation for DT:

### Cross Validation for DT

```
In [77]: for i in range(2,6):  
         DT_Val=cross_val_score(DT,X_Scaled,Y_new,cv=i)  
         print("The cross validation score for Decision Tree Regressor",i,"is",DT_Val.mean())
```

The cross validation score for Decision Tree Regressor 2 is 0.7001601206866955  
The cross validation score for Decision Tree Regressor 3 is 0.7354424014889114  
The cross validation score for Decision Tree Regressor 4 is 0.7002785021394589  
The cross validation score for Decision Tree Regressor 5 is 0.7238830226633229

## Cross Validation for RF:

### Cross Validation for RF

```
In [79]: for i in range(2,6):  
         RF_Val=cross_val_score(rf,X_Scaled,Y_new,cv=i)  
         print("The cross validation score for",i,"is",RF_Val.mean()*100)
```

The cross validation score for 2 is 83.45543872118004  
The cross validation score for 3 is 85.87326337278839  
The cross validation score for 4 is 86.04839891223044  
The cross validation score for 5 is 85.63141513051768

## Cross Validation for GB:

### Cross Validation for GB

```
In [80]: for i in range(2,6):  
         GB_Val=cross_val_score(GB,X_Scaled,Y_new,cv=i)  
         print("The cross validation score for",i,"is",GB_Val.mean()*100)
```

The cross validation score for 2 is 85.7902333179118  
The cross validation score for 3 is 87.15819308502503  
The cross validation score for 4 is 86.93685769382368  
The cross validation score for 5 is 86.57363515177636

***As none of the models are overfitted, and based on the  $r$  squared score and cross validation scores, Gradient Boosting Regressor model is best for this dataset.***

### **Hypertuning Parameter:**

As gradient boosting has been termed as the best model for this dataset. Lets try to tune the parameters to see if the score can be increased.

#### **1) Trying first with best parameter by Grid Search CV method:**

```
In [82]: gb=GradientBoostingRegressor()

In [84]: Parameters={'criterion':['mse', 'mae'], 'min_samples_split':[3,4], 'min_samples_leaf':[2,3]}
         clf=GridSearchCV(gb,Parameters)
         clf.fit(X_train,y_train)

Out[84]: GridSearchCV(estimator=GradientBoostingRegressor(),
                      param_grid={'criterion': ['mse', 'mae'],
                                   'min_samples_leaf': [2, 3],
                                   'min_samples_split': [3, 4]})

In [85]: clf.best_params_ #taking the best parameters

Out[85]: {'criterion': 'mse', 'min_samples_leaf': 3, 'min_samples_split': 4}
```

Now reinitialising the best parameters and checking the score:

```
In [86]: gb=GradientBoostingRegressor(criterion='mse',min_samples_leaf=3,min_samples_split=4)

In [87]: X_train,X_test,y_train,y_test=train_test_split(X_Scaled,Y_new,test_size=0.20,random_state=50)
         gb.fit(X_train,y_train)
         pred_test=gb.predict(X_test)

         print('R-Squared:',r2_score(y_test,pred_test)*100)

R-Squared: 86.93172876007162
```

Score has increased, but not much hence trying the next method to see the parameters can be tweaked further thereby increasing the score.

## 2) *Trying different parameters other than Grid Search CV.*

```
In [90]: gb=GradientBoostingRegressor(criterion='mae',min_samples_leaf=2,min_samples_split=3)
```

```
In [91]: X_train,X_test,y_train,y_test=train_test_split(X_Scaled,Y_new,test_size=0.20,random_state=50)
         gb.fit(X_train,y_train)
         pred_test=gb.predict(X_test)

         print('R-Squared:',r2_score(y_test,pred_test)*100)
```

R-Squared: 88.24863201774133

**Score increased, thereby saving the model**

# Conclusion

### **Key Findings:**

Going through the data, it can be concluded that the following are the main 25 parameters that influence the price of houses in Australia. With the help of these details Surprise Housing will be able to manipulate the price and enter the Australian market.

1. Overall Qual
2. Above grade (ground) living area square feet
3. Garage Cars
4. Basement Quality
5. Basement Condition
6. External Quality
7. GarageArea
8. Total Basement Square feet
9. Kitchen Quality
10. First Floor Squarefeet
11. Full Bath
12. Garage Finish
13. Total rooms above grade (does not include bathrooms)
14. Year on which House wasBuilt
15. Year on which renovation was added
16. Masonry veneer area in square feet
17. Fire places
18. Year on which garage was built
19. Heating QC
20. Foundation
21. Type 1 finished square feet
22. Open Porch square feet
23. 2<sup>nd</sup> Flr SF
24. Lot Frontage
25. WoodDeck SF